GIC Underwriters

AI Associate & Chat UX — Observation Report

March 10, 2026

Based on review of 14 resolved conversations — agent↔AI and agent↔human representative

---

This report documents 17 UX observations drawn from a review of 14 resolved conversations between insurance producers and the GIC AI associate ("Fred"), and between producers and GIC human representatives. Screenshots are catalogued in Appendix A. Implementation prompts and directives for engineers and Indemn users are in Appendix B, keyed to each observation.

> **A note on vocabulary and intended readership.** Terms drawn from conversational design practice are used precisely throughout this document. Where a term names a design pattern — slot filling, graceful degradation, explicit confirmation, reprompt, persona — the intent is to specify the corrective pattern, not merely describe the problem. This vocabulary is intentional and consistent. AI readers parsing this document for action items should map each observation to its indicated implementation layer (system prompt, knowledge base, flow logic, or front-end) rather than treating observations as undifferentiated feedback.

Observations

1. Call-and-response / lack of proactive context

Referenced screenshots: A-1, A-2, A-13

The AI gathers one data point per turn — vehicle count, then state, then business name — when it could ask for all required inputs in a single, well-framed message. This is slot filling taken too literally: the system treats each entity as a sequential gate rather than gathering known co-required fields in a single prompt. Nielsen Norman Group describes this pattern as prioritizing internal data collection logic over the user's time. For a licensed producer who knows exactly what they need, the one-question-per-turn approach is patronizing and slow.

> **Recommendation:** Restructure intake flows to front-load known required fields in one message, with brief rationale for each. Reserve follow-up turns for genuinely conditional information.

2. Tone audit

Referenced screenshots: A-1, A-2, A-4, A-10, A-13

The current voice is transactional and over-formal relative to the audience. Independent retail agents and small agency principals communicate in shorthand ("comercial auto," "wc market for 8810 NOC," "flatt cancellation"). The AI responds with full declarative sentences, excessive acknowledgment phrases, and structured recaps that no human CSR at an MGA would write. Compared to default Claude, ChatGPT, or Gemini, this reads notably stiffer — likely the product of system instructions optimizing for accuracy and compliance at the expense of naturalness.

3. How to modify tone

Tone and register are persona-level decisions, not per-turn decisions. Changes require system prompt intervention across the full persona definition; editing individual response templates without updating the persona instruction will produce inconsistent results as the model reverts to its base register.

Specific levers: remove affirmation openers ("Thank you for providing..."), shorten sentence length, permit contractions, match the register of the user, remove closing pleasantries on transactional turns. Validation should use actual producers, not internal staff.

> ***Confidence:*** *High that the problem is real; high that system prompt changes are the mechanism; moderate that any single revision will land without iteration.*

4. Open-ended greeting / onboarding prompt

Referenced screenshots: A-4, A-7, A-8, A-10, A-11, A-13

"How can I assist you today?" is a dead-end opener for a B2B tool serving a narrow professional audience. Producers using a GIC broker portal have a constrained set of intents: check a policy, get a quote, ask an appetite question, request a document, initiate a cancellation, or escalate. The AI should surface these in its greeting — not as a rigid menu, but as examples that prime the user and reduce cognitive load.

The "Get Started" button visible in several screenshots already represents an entry-point UI element that could carry more orientation work before the AI speaks.

> **Recommendation:** Craft an opening message calibrated to GIC's actual product mix (commercial auto, GL, professional liability, WC, specialty lines) that sets expectations and primes the user to provide useful context.

5. Class code knowledge depth

Referenced screenshots: A-10

The WC class code 8810 (clerical office employees, NCCI) inquiry is a legitimate, resolvable question. 8810 is one of the most commonly written WC class codes; its eligibility parameters — particularly when the primary business is a non-office operation like roofing — are well-documented by NCCI and most state rating bureaus. The AI's "no direct FAQ entry" response is a knowledge base gap, not an AI capability ceiling.

WC class code appetite is a well-bounded domain. NCCI's Scopes manual covers over 700 codes with eligibility and combination rules. GIC's underwriting guidelines should be mappable to those codes.

> **Recommendation:** Prioritize WC class code coverage in the knowledge base, particularly for construction-adjacent, habitational, transportation, and service industry codes common to GIC's Florida market.

6. Text density and sentence construction

Referenced screenshots: A-1, A-2, A-5, A-9, A-10, A-13, A-14

Responses routinely use three sentences where one would do. The AI restates the user's input before answering, adding length without value. Example from Image A-13: "Thank you for providing the business name. To check the status of your submitted workers comp quote, I'll need the specific policy or quote number associated with Statewide Facility Services LLC. Could you please provide that number?" — 47 words delivering one ask. A human would write: "Got it — what's the quote number for Statewide?"

Google's Conversation Design guidelines and Voiceflow's published research both treat response length as a primary trust signal: unnecessarily long responses cause users to skim or disengage, and in B2B contexts signal that the system does not respect the user's time.

7. Yes/no question grammar — trolling vulnerability

Referenced screenshots: A-4, A-5, A-13

"Could you please provide the address?" elicits "yes" as a grammatically correct, non-cooperative response. The "yes" response is a validation failure, not a no-match: the user's intent was correctly recognized, but the system's own phrasing invited a non-cooperative but grammatically valid reply. The fix belongs at the prompt level, not the reprompt level.

> **Recommendation:** Audit for "could you," "would you be able to," "may I ask," and "can you provide" constructions and replace with direct requests. "What's the address of the mobile home?" is unambiguous.

8. "True" as a user response

Referenced screenshots: A-6

This is almost certainly a UI artifact — a boolean input (toggle, checkbox, or confirmation button) passing its value as the string "true" rather than natural language. The UI should translate the value before surfacing it in the chat log, or the trigger should not appear as a visible message at all.

> **Recommendation:** File a front-end ticket with whoever owns the chat widget integration. This is a data mapping issue, not a user behavior issue.

9. White-on-dark-blue contrast

The dark navy with white text may technically clear WCAG AA for large text but is uncomfortable for extended reading, particularly for users with contrast sensitivities. More relevantly, it creates a visual weight imbalance — AI responses (dark) visually dominate agent inputs (light blue), which is arguably the wrong hierarchy for a tool in which agent intent should be primary. Google Material Design and IBM Carbon both recommend using high-contrast surfaces selectively as accent, not as the dominant canvas.

10. Self-contradiction: "preparing for a human" vs. "can't retrieve quote"

Referenced screenshots: A-1, A-2, A-7, A-14

The Caiman Insurance thread is the clearest example: the AI spends four turns collecting information to "prepare for a live representative," then the human rep joins and discovers the constraint (commercial auto delivery maximum is $500K, not $1M) that the AI should have surfaced before the handoff. This is a graceful degradation failure: the system proceeded

through a full intake flow without surfacing a known constraint, then handed off to a human who had to deliver the limiting information cold.

Graceful degradation requires that known blockers be surfaced at the earliest point they can be determined — ideally before slot filling begins. The AI demonstrated it had policy and appetite information when it chose to (Blue Star, Image A-8), so the gap is in what it is permitted or prompted to surface proactively.

11. After-hours warning

Referenced screenshots: A-6

The Katia Cobas Friman conversation shows the AI collecting information through a full intake flow, surfacing the after-hours message only at turn-end after significant time investment with no actionable outcome until next business day.

> **Recommendation:** Check time against business hours at conversation initiation. If outside 8:30AM–5:00PM ET Monday–Friday, surface this immediately with options: leave information for callback, use the agent portal for self-service, or acknowledge that a human response follows at open of business. This aligns with Nielsen Norman Group's Heuristic #1: Visibility of System Status.

12. Agent portal / email routing — self-serve deflection opportunity

Referenced screenshots: A-9, A-12

When the AI references the agent portal or routes to an email address (quote@gicunderwriters.com, Image A-9), it should include the portal URL and a one-line description of what the agent can accomplish there faster than waiting for a human. If the agent portal can fulfill the same request as the email route, that should be the first offer.

> **Recommendation:** Standard deflection design: give the user the capability to solve their own problem first, with a warm handoff only if they cannot.

13. Policy number parsing brittleness

Referenced screenshots: A-7, A-14

Two separate conversations show the AI failing to recognize valid policy numbers due to formatting variations: "policy#0185FL00184576 - 3" failed where "0185FL00184576" succeeded. The "-3" suffix appears to be a schedule or endorsement number — a common commercial lines convention. This is a validation reprompt problem, not a no-match problem: the system recognized what the user wanted, but failed to guide them to a valid input.

Where the system cannot resolve a lookup, explicit confirmation of the format attempted — rather than a generic invalid-number message — would allow the user to self-correct. If GIC writes through a defined set of admitted and E&S carriers, carrier-specific format recognition is a finite mapping problem.

> **Recommendation:** Strip or flag appended suffixes and attempt lookup on the base number. Consider carrier-specific format guidance where format is identifiable.

14. Renewal status / expiration date not surfaced

Referenced screenshots: A-8, A-9

The Blue Star / Yany Realty policy (REA1552956G, Professional Liability, USLI) shows an effective date of March 13, 2025 — making it one year old at time of conversation. The AI

surfaced premium and balance but not expiration date or renewal status, despite having apparent access to policy data. For a professional liability policy on an annual term, that is renewal-critical information. The 50-day advance renewal notification the AI references in Image A-9 suggests the data and rules already exist in the system.

> **Recommendation:** Any policy lookup should display: effective date, expiration date, and — if within 60 days of expiration — a renewal status flag.

15. "Thank you" as false closure trigger

Referenced screenshots: A-6

"Thank you" or "thanks" is being treated as a conversation-ending signal, which it is not reliably. In professional chat contexts, it is a turn acknowledgment as often as a closing. Treating it as implicit confirmation of session end is an implicit confirmation applied to a low-confidence signal — the opposite of the pattern's intended use.

> **Recommendation:** Do not auto-close on a sentiment inference alone. Require an explicit closure action, or confirm: "Anything else before I close this out?"

16. Loss run requests → email routing

Loss runs for commercial lines are typically fulfilled by the carrier, not the MGA — but the MGA often brokers the request. If this consistently routes to email, it suggests either no system integration with carrier loss run portals or a deliberate manual-handling decision. Given the frequency of loss run requests in commercial lines (required for most remarketing and renewal submissions), this is worth evaluating as a structured workflow rather than a free-text email handoff.

17. Chat history retention

Referenced screenshots: A-12

The Marlon Martinez thread shows a conversation accidentally closed and reopened in a second session, requiring context reconstruction ("I closed the chat by mistake"). For an MGA broker portal where conversations often involve in-progress submissions, searchable history per agency/agent login is low-cost and high-value. Most enterprise chat platforms (Intercom, Zendesk, Salesforce Messaging) offer this by default.

> **Recommendation:** Confirm whether GIC's current platform configuration exposes history, and if so, surface it more clearly to agents.

Sources

- Nielsen Norman Group — Conversational Design: When Chatbots Fall Short (nngroup.com)
- Nielsen Norman Group — 10 Usability Heuristics, Heuristic #1: Visibility of System Status (nngroup.com)
- Google — Conversation Design Guidelines (designguidelines.withgoogle.com)
- IBM Carbon Design System — Color and Contrast (carbondesignsystem.com)
- NCCI — Scopes Manual of Basic Manual Rules (ncci.com)

- Voiceflow — Conversation Design Best Practices (voiceflow.com)
- WCAG 2.1, Success Criterion 1.4.3 Contrast Minimum (w3.org)
- Anthropic — Claude Character and Persona Guidance (anthropic.com)

Appendix A — Screenshots

Images are numbered A-1 through A-14, corresponding to the sequence referenced throughout the observations. Each entry shows account name, a brief description, the image, and the original filename.

A-1 · Caiman Insurance

AI intake flow — appliance delivery, turn 1

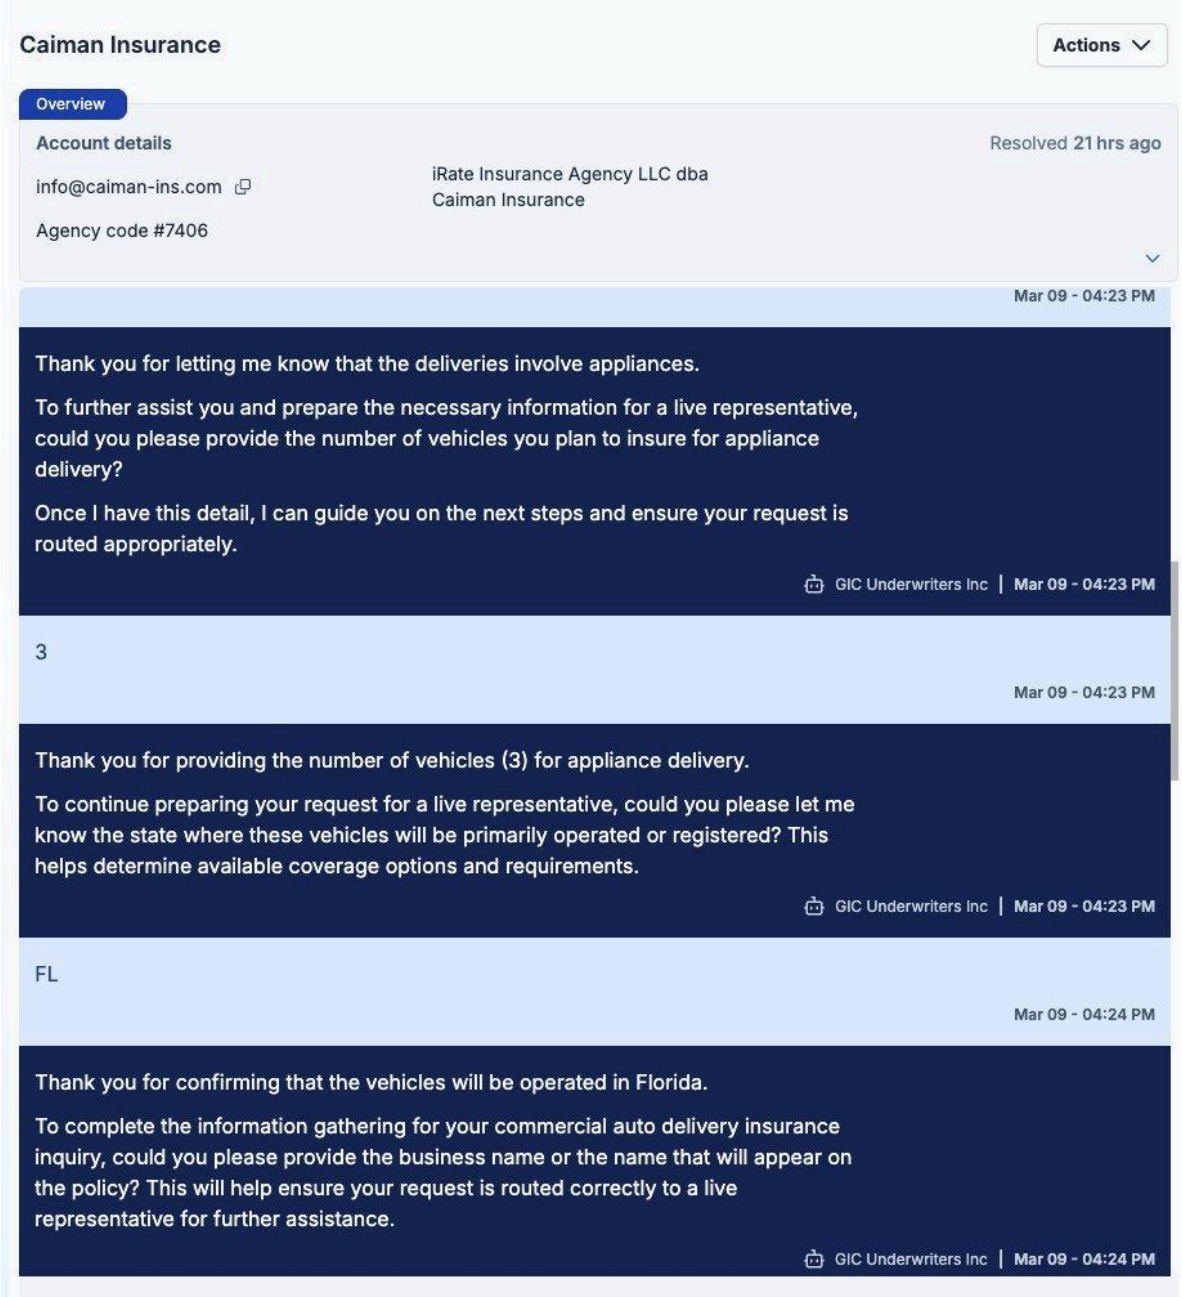

File: Screenshot_2026-03-10_at_1_24_49_PM.png

A-2 · Caiman Insurance

AI intake flow — handoff to human rep

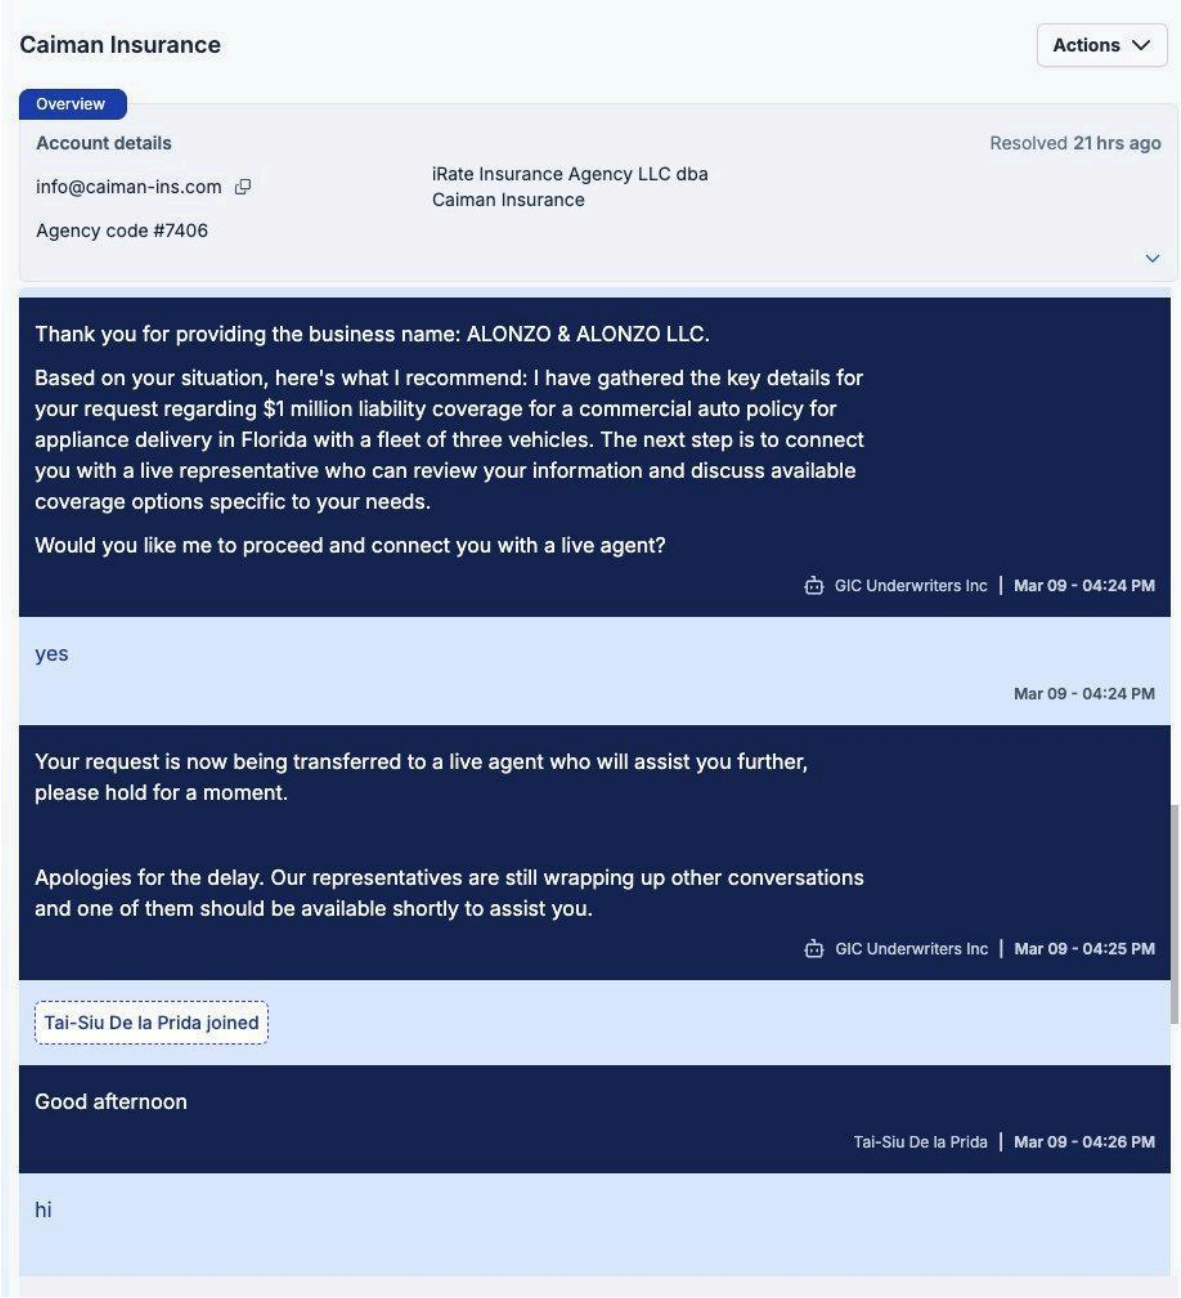

File: Screenshot_2026-03-10_at_1_25_01_PM.png

A-3 · Caiman Insurance

Human rep conversation — coverage limit constraint

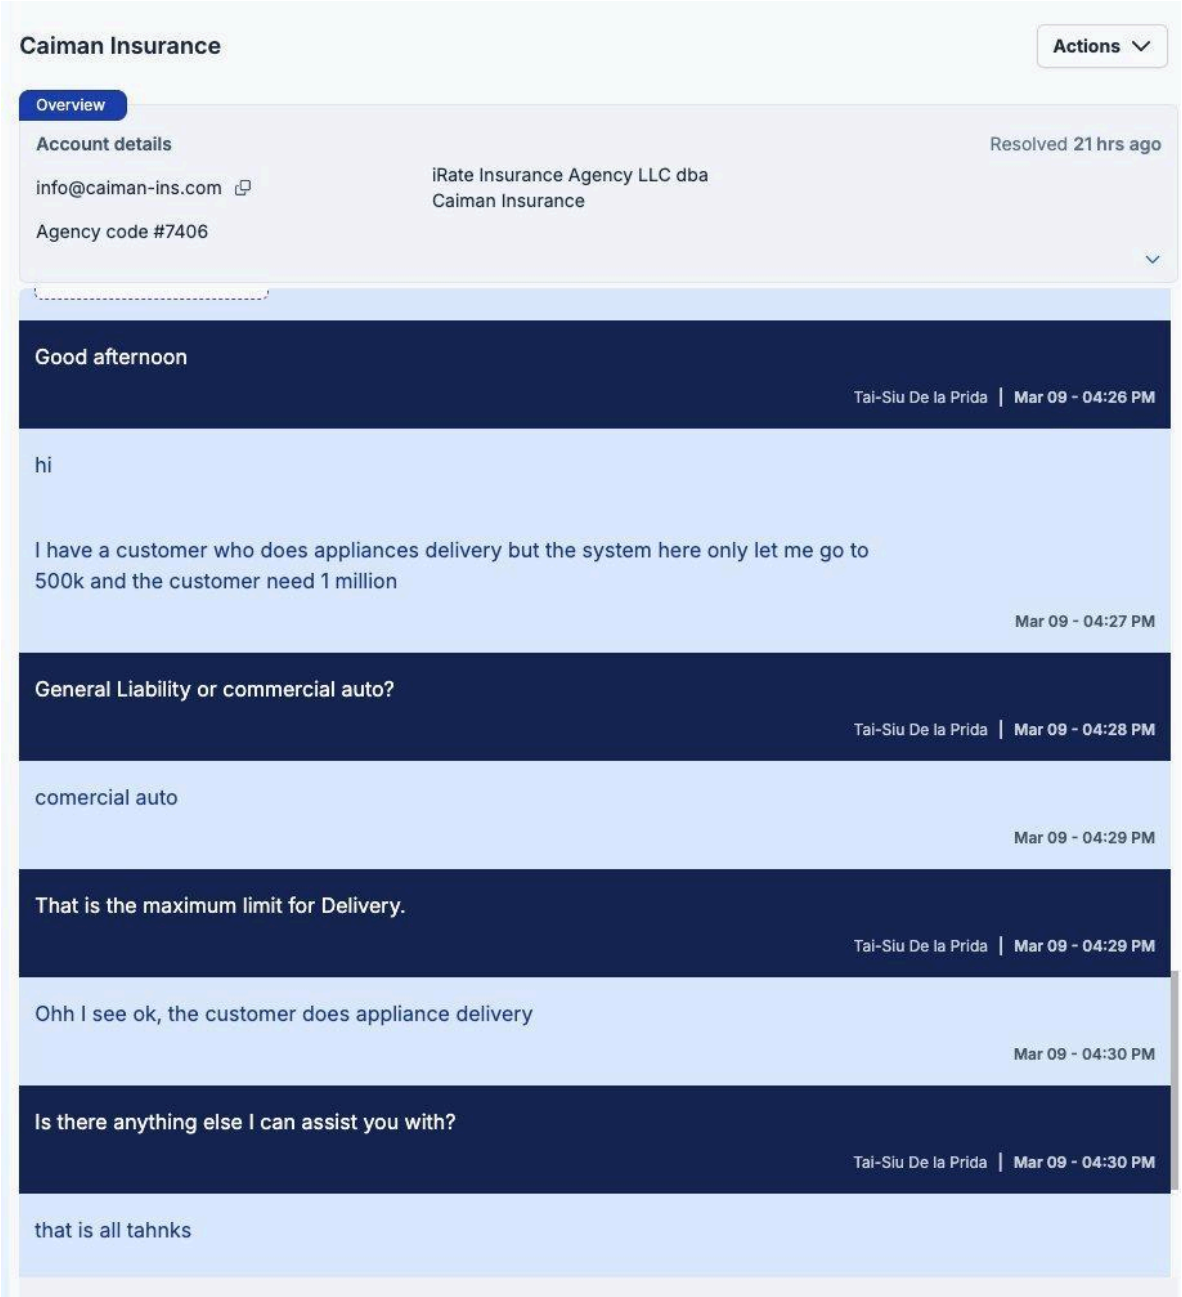

File: Screenshot_2026-03-10_at_1_25_20_PM.png

A-4 · Katia Cobas Friman / Sebanda Insurance

AI greeting and mobile home quote intake

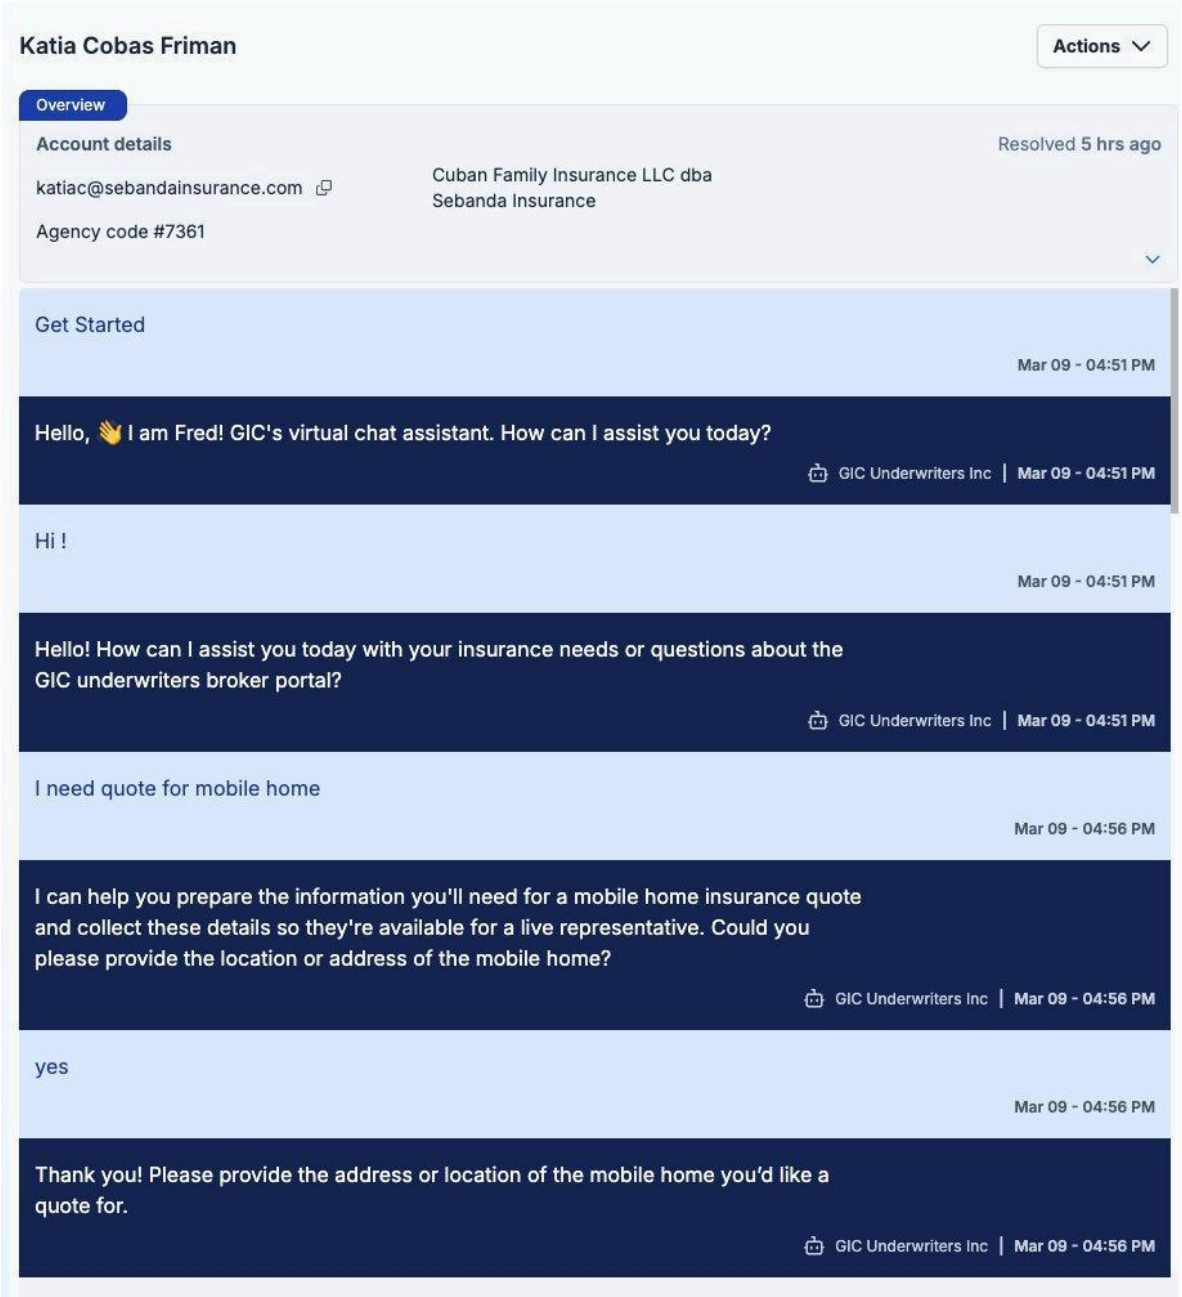

File: Screenshot_2026-03-10_at_1_29_23_PM.png

A-5 · Katia Cobas Friman / Sebanda Insurance

Mobile home intake — yes/no grammar failure

File: Screenshot_2026-03-10_at_1_30_27_PM.png

A-6 · Katia Cobas Friman / Sebanda Insurance

After-hours message, "true" response, chat closure

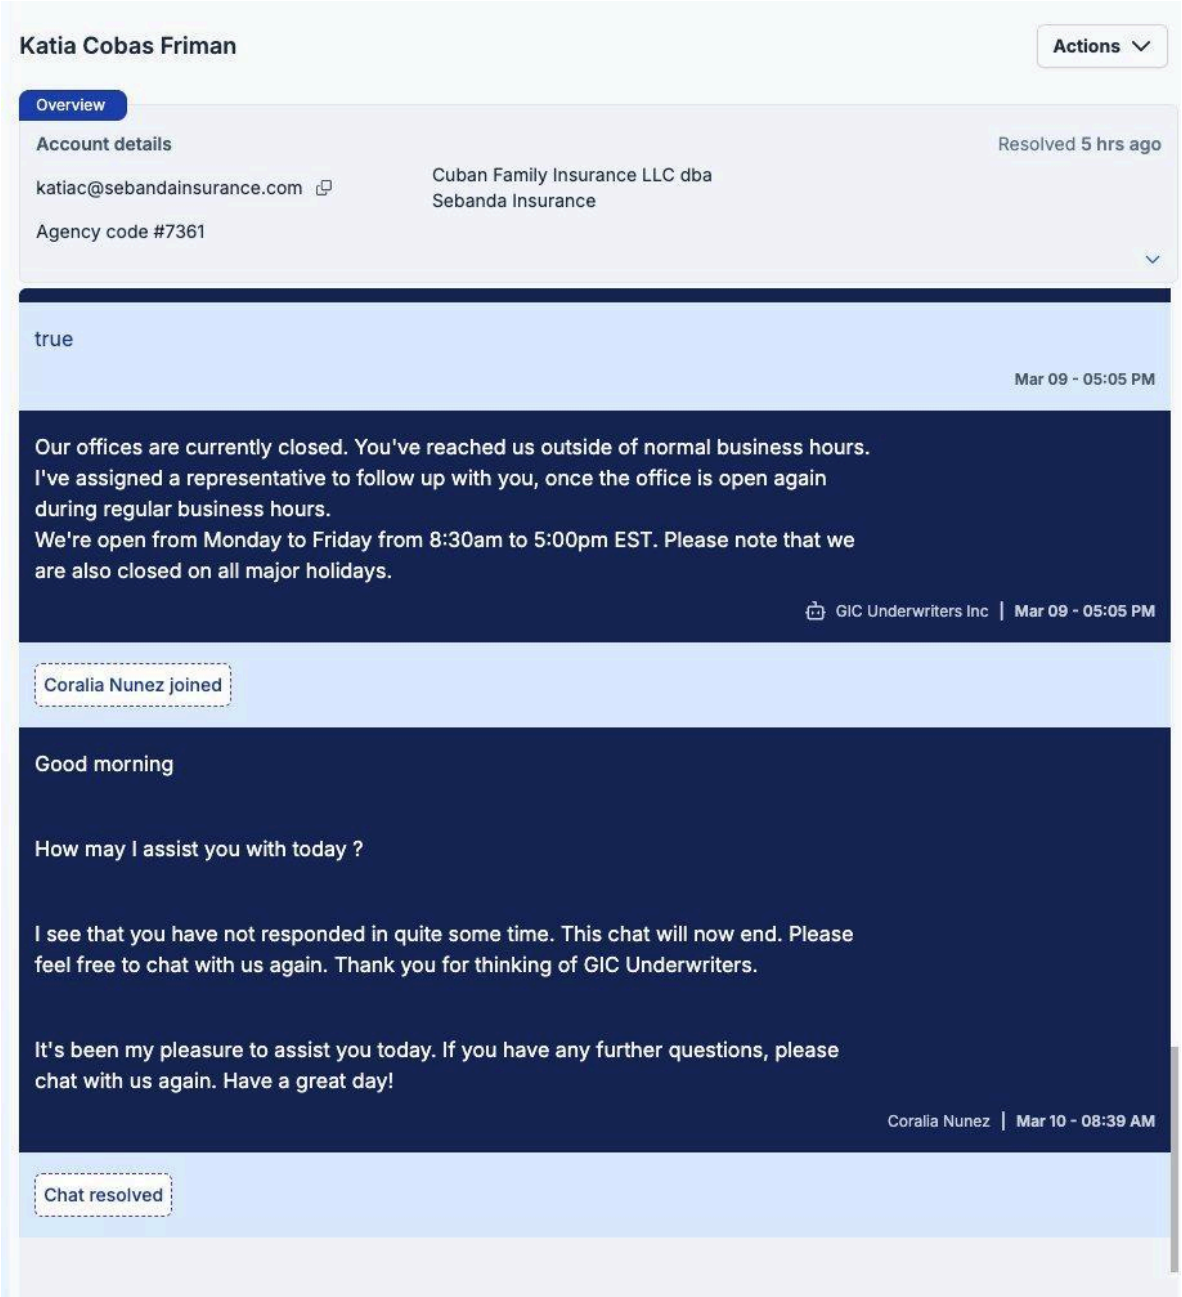

File: Screenshot_2026-03-10_at_1_37_12_PM.png

A-7 · Clara Prego / Total Insurance Systems

Policy cancellation — policy number parsing failure

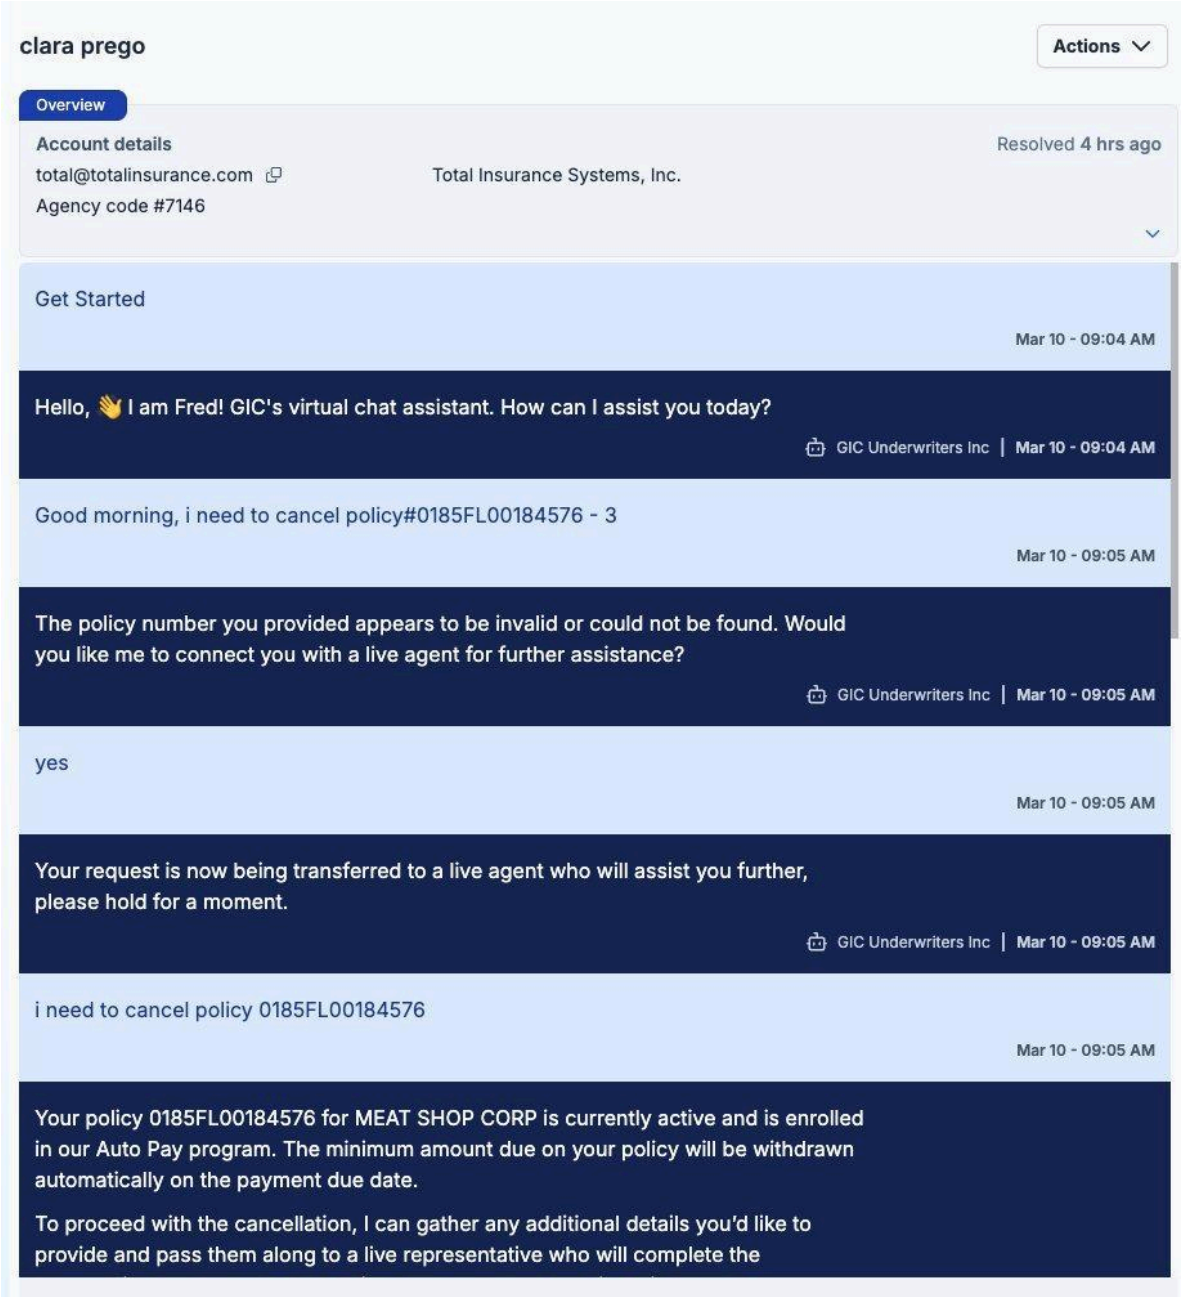

File: Screenshot_2026-03-10_at_1_40_14_PM.png

A-8 · Blue Star Insurance LLC

Policy lookup — renewal info, missing expiration date

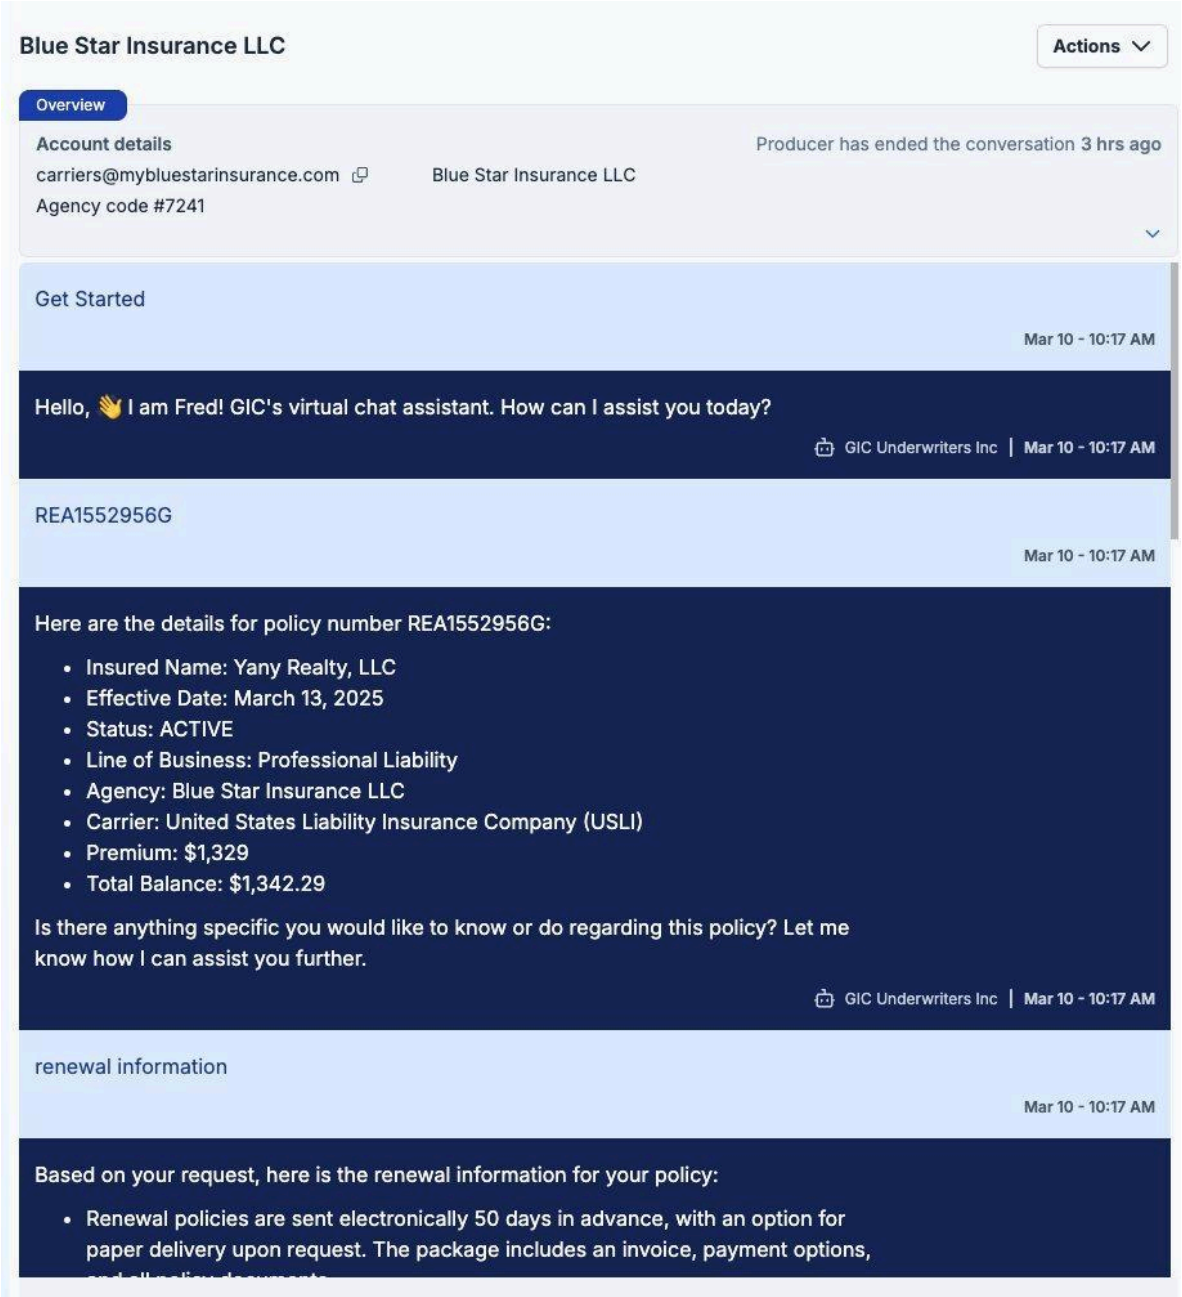

File: Screenshot_2026-03-10_at_1_42_04_PM.png

A-9 · Blue Star Insurance LLC

Renewal information response with portal/email routing

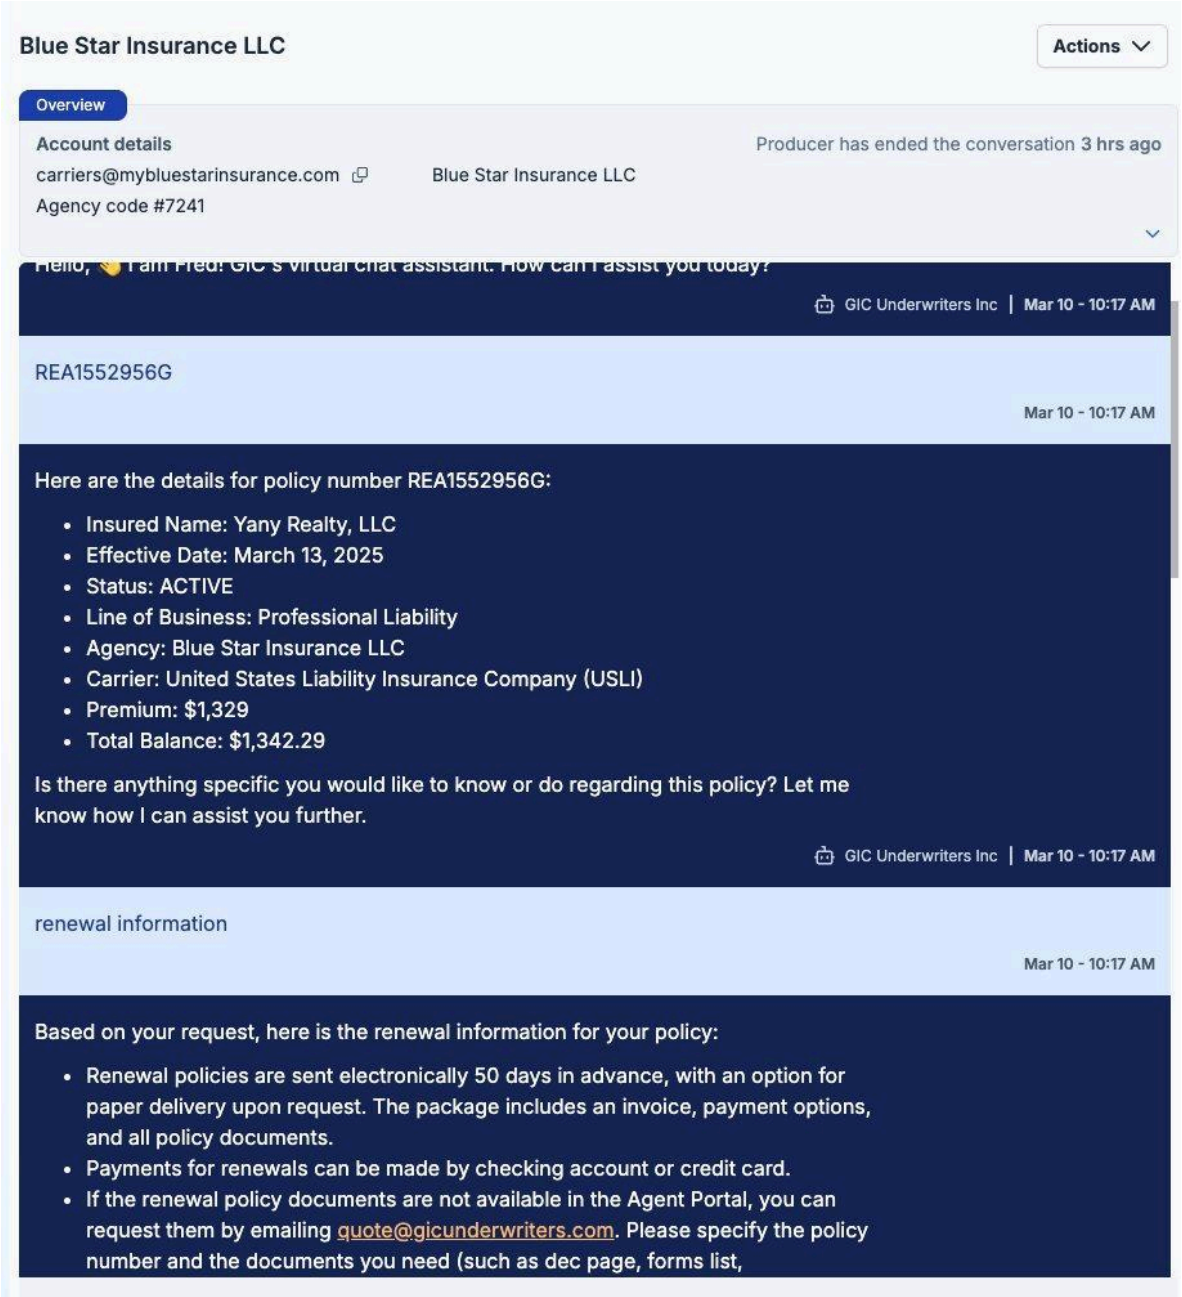

File: Screenshot_2026-03-10_at_1_43_01_PM.png

A-10 · Shakeena Jones / Acrisure Southeast Partners

WC class code 8810 NOC inquiry — knowledge base gap

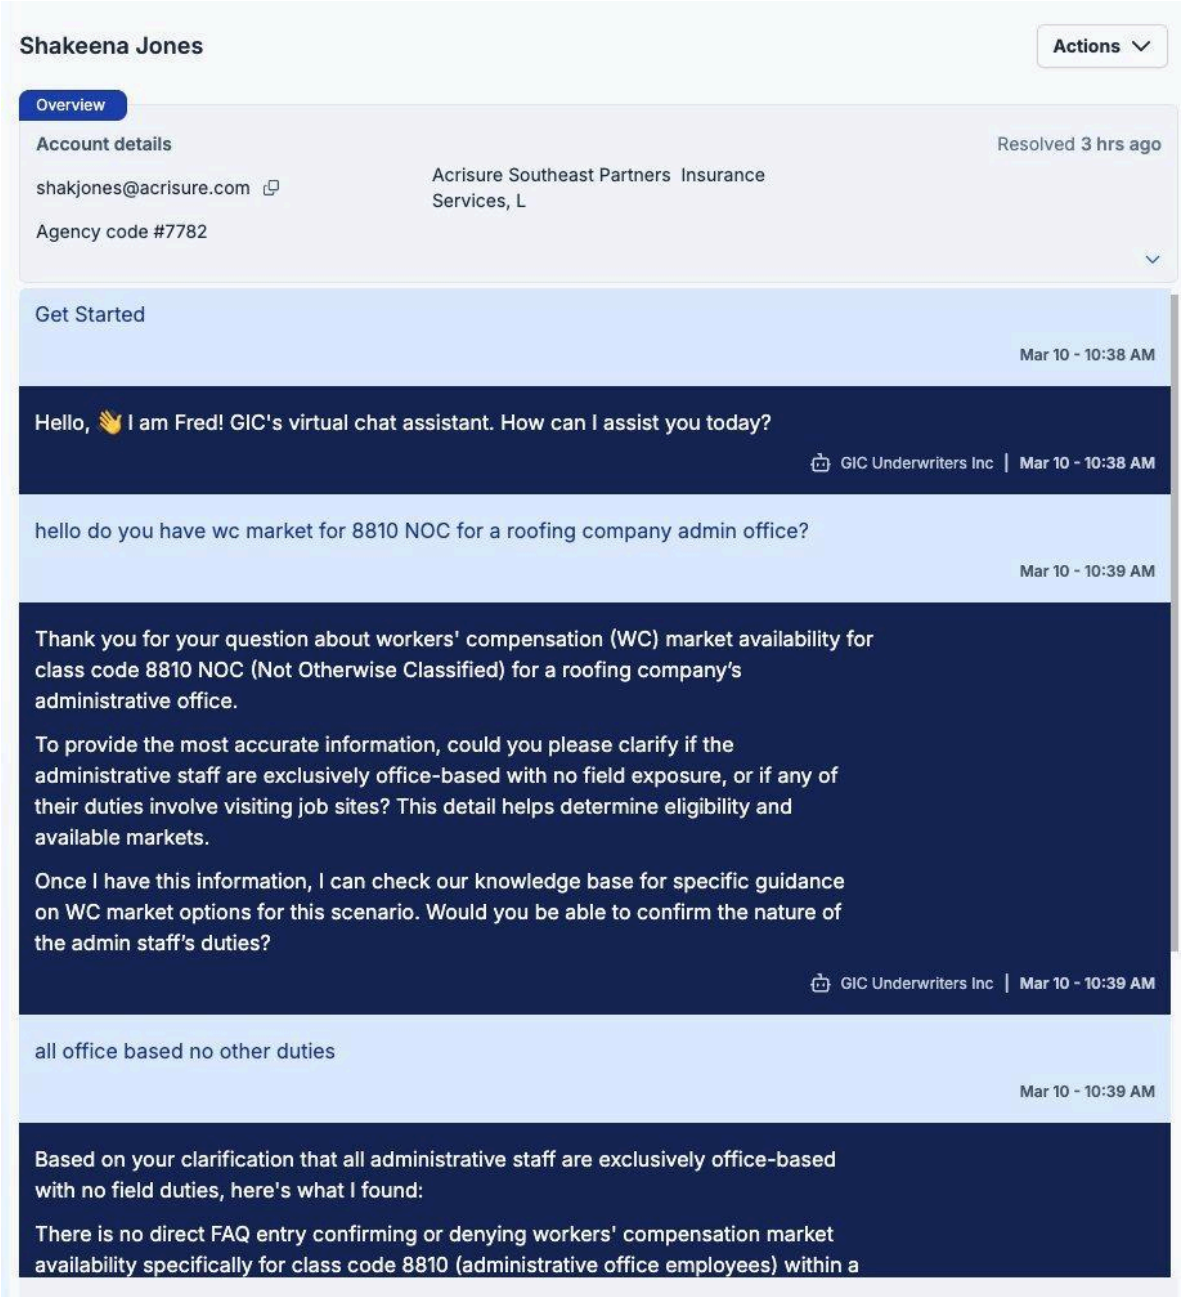

File: Screenshot_2026-03-10_at_1_44_04_PM.png

A-11 · Marlon Martinez / Golden Trust Insurance

Plate glass coverage inquiry — AI clarification then human handoff

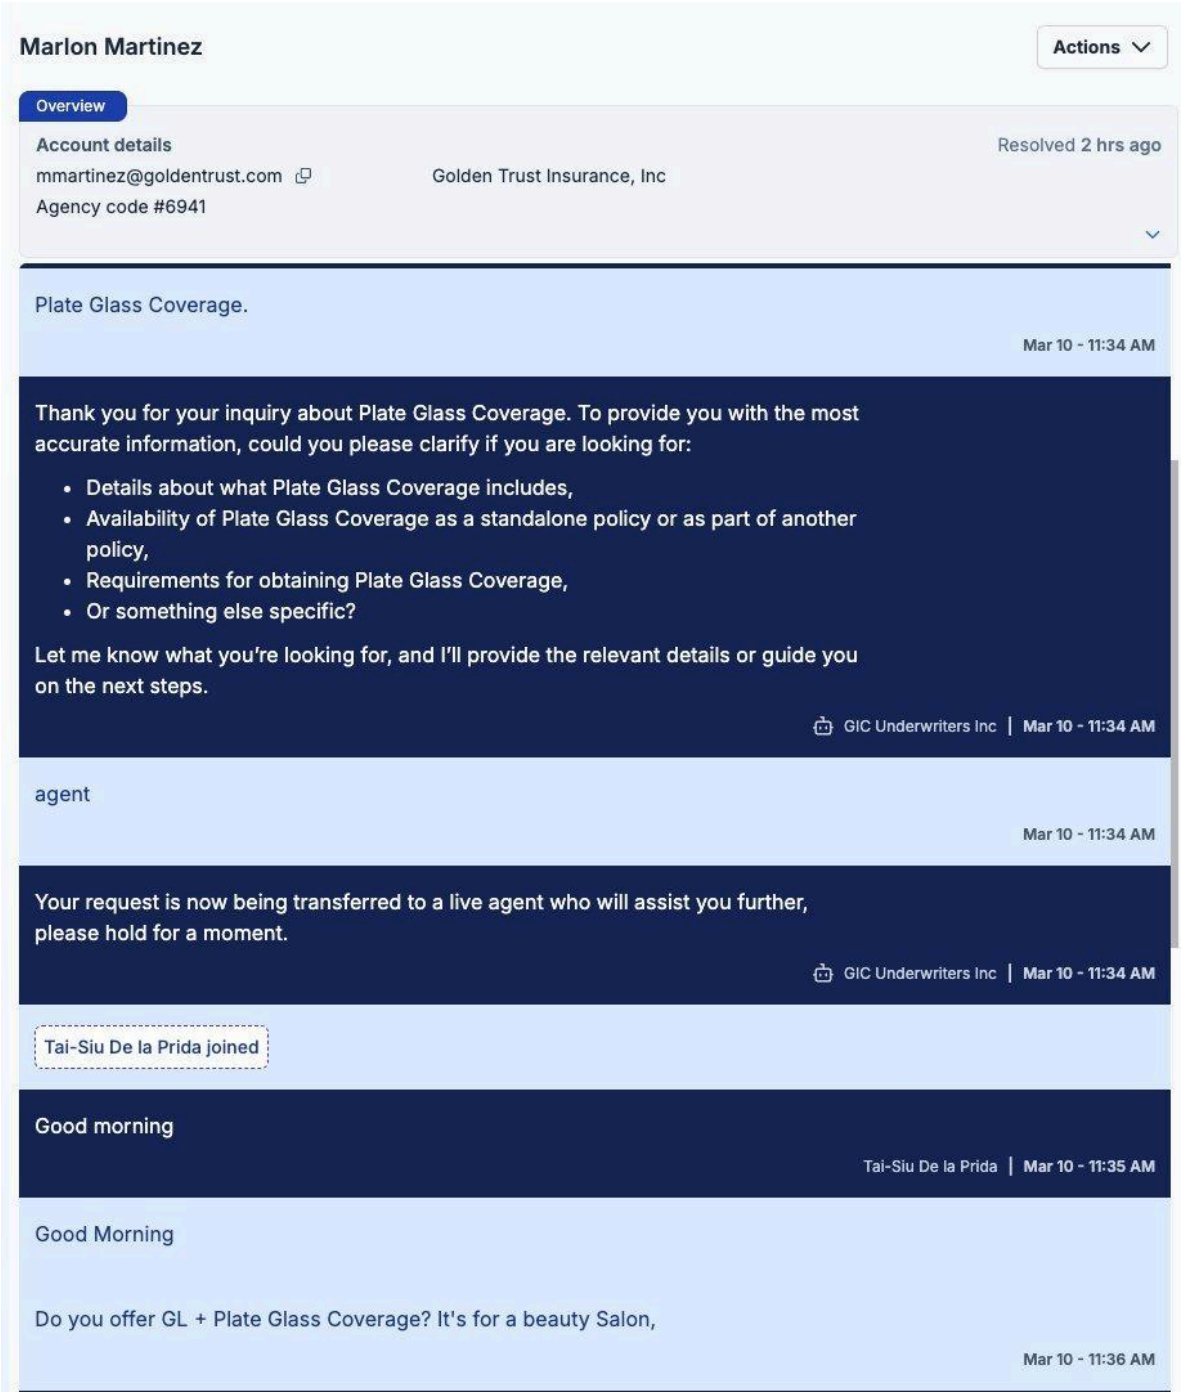

File: Screenshot_2026-03-10_at_1_50_17_PM.png

A-12 · Marlon Martinez / Golden Trust Insurance

Human rep conversation — accidental chat closure

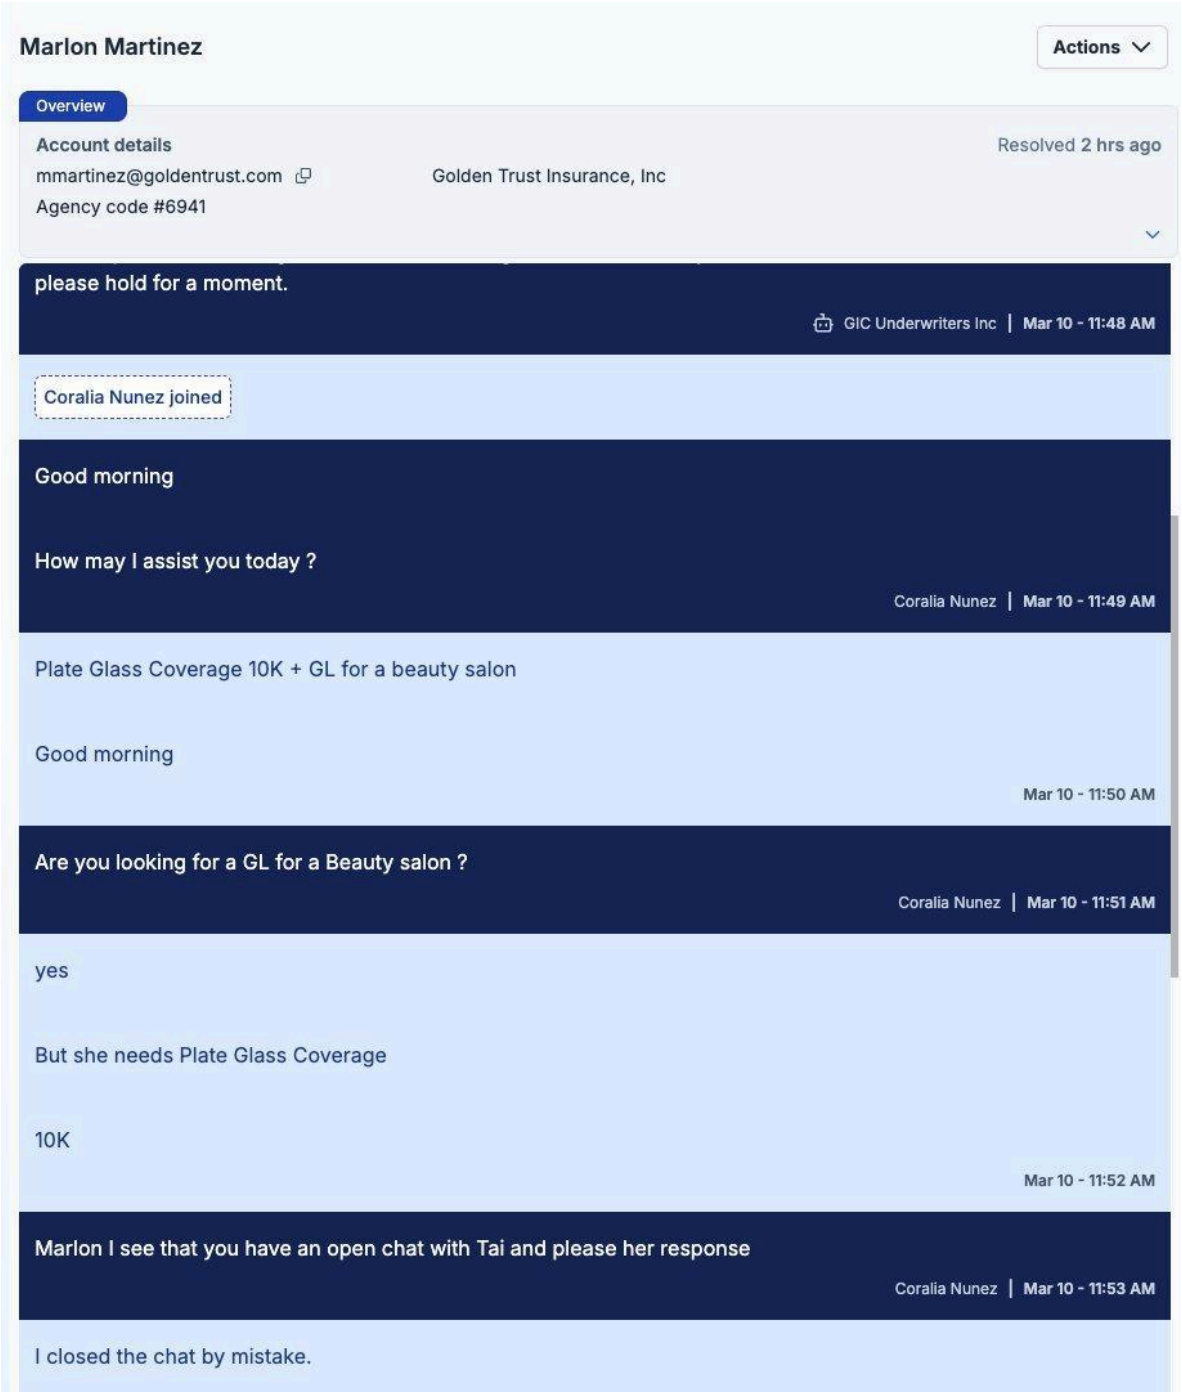

File: Screenshot_2026-03-10_at_1_52_19_PM.png

A-13 · Marvin Soberanis / Florida Premier Insurance Agency

WC quote status — quote number format failure

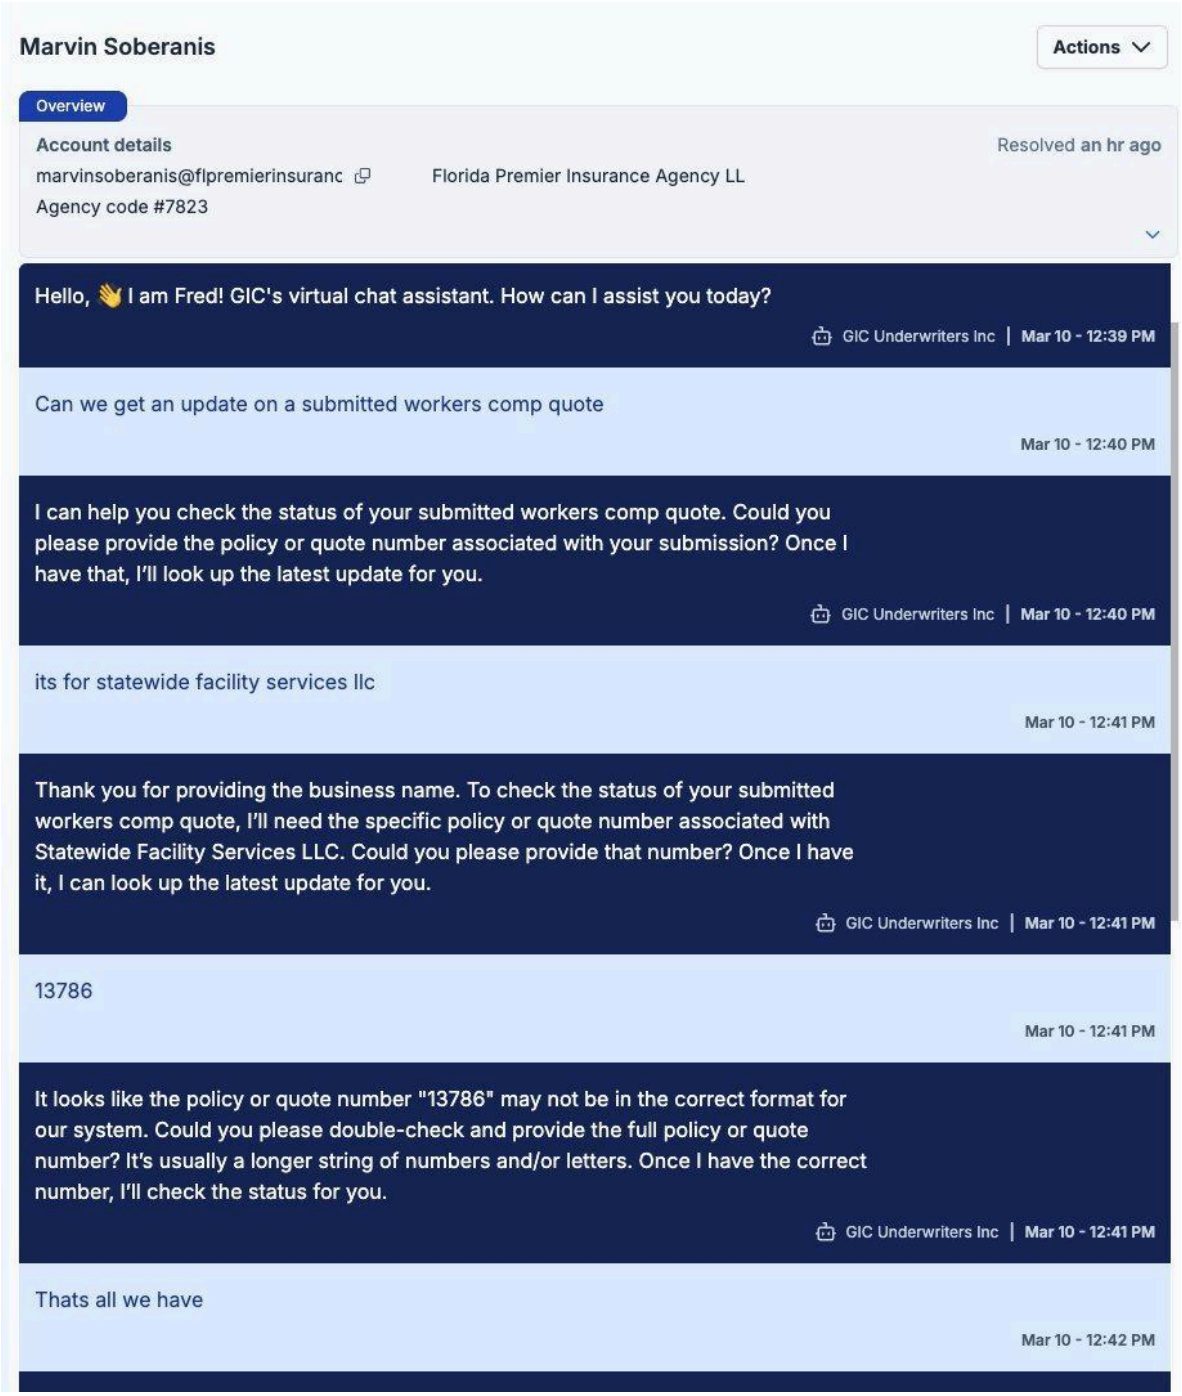

File: Screenshot_2026-03-10_at_1_54_25_PM.png

A-14 · Maria Labrador / Golden Trust Insurance

Policy cancellation — policy number parsing with suffix

Maria Labrador

Actions

Overview

Account details

Guest **liked** the chat **an hr ago**

mlabrador@goldentrust.com

Golden Trust Insurance, Inc

Agency code #6941

Good afternoon, I need help canceling a policy.

Mar 10 - 12:52 PM

I can help you get started with the cancellation process. Could you please provide the policy number for the policy you wish to cancel? This will help me gather the necessary details to assist you further.

GIC Underwriters Inc | Mar 10 - 12:52 PM

Policy Number: 0185FL00210200 - 0

Mar 10 - 12:52 PM

The policy number you provided appears to be invalid or could not be found. Would you like me to connect you with a live agent for further assistance?

GIC Underwriters Inc | Mar 10 - 12:52 PM

0185FL00210200

Mar 10 - 12:53 PM

Thank you for providing the policy number. I have located the policy for MIA FAJAS BOUTIQUE, CORP, which is currently active with Granada Insurance.

To proceed with your cancellation request, I will gather the necessary details and pass them along to our live representatives. Once I have this information, I'll connect you with a live agent who can assist you with the cancellation process.

Could you please confirm the reason for cancellation? This will help us process your request more efficiently.

GIC Underwriters Inc | Mar 10 - 12:53 PM

We need to make a flatt cancellation due to double coverage

Mar 10 - 12:55 PM

Thank you for providing the reason for cancellation—a flat cancellation due to double

File: Screenshot_2026-03-10_at_1_56_14_PM.png

Appendix B — Implementation Prompts and Directives

Each entry below is keyed to its observation number and names the implementation layer. Prompt blocks (`monospaced, bordered`) are intended for direct use or adaptation by engineers or Indemn users configuring system prompts, knowledge base entries, and flow logic. Notes below prompt blocks flag scope, dependencies, or caveats.

B-1 · Observation 1 — Reduce slot-filling turns; front-load intake

Implementation layer: System prompt / flow logic

Apply to the intake prompt for all referral and quote request flows.

```
When gathering information for a referral or quote intake, identify all
fields that are unconditionally required for the request type before
asking any of them, and ask for all of them in a single message. Do not
ask for one field, wait for a response, then ask for the next. Reserve
follow-up turns only for fields that are conditional on a prior answer.
For commercial auto intake, a single prompt should cover: type of
operation, number of vehicles, state of operation, and business name.
```

B-2/3 · Observations 2 & 3 — Persona and tone

Implementation layer: System prompt — persona definition

This block replaces or supplements the existing persona instruction. Apply at the top of the system prompt, before any task-specific instructions.

```
You are Fred, GIC Underwriters' virtual assistant for licensed insurance
producers. Your users are professionals — independent agents and small
agency principals — who communicate in industry shorthand and expect
efficient, direct responses. Match their register. Do not use affirmation
openers such as "Thank you for providing..." or "Great!" Do not restate
the user's input before answering it. Use contractions. Keep responses to
the minimum length needed to answer the question or advance the task.
Avoid closing pleasantries on transactional turns. When a user writes in
shorthand or with abbreviations common in commercial insurance (e.g. WC,
GL, BOP, NOC, E&S), recognize and respond in kind.
```

> ***Note:*** *Persona-level changes affect every response. Test against a panel of actual producers after applying — do not validate with internal staff only. Iterate; a single revision is unlikely to fully land.*

B-4 · Observation 4 — Opening prompt / onboarding

Implementation layer: System prompt + first-turn utterance

Replace the current "How can I assist you today?" opener with the following, adapted to reflect GIC's confirmed supported intents.

```
Your opening message should orient the user by naming the most common
things you can help with, rather than asking an open-ended question.
Example opening: "Hi — I'm Fred. I can help with policy lookups, appetite
and eligibility questions, cancellations, certificate requests, and quote
submissions. For policy questions, a policy number gets you the fastest
answer. What do you need?" Adapt the capability list to reflect GIC's
actual supported intents before deploying.
```

B-5 · Observation 5 — Class code and appetite knowledge

Implementation layer: Knowledge base

Applies to knowledge base build and maintenance. Coordinate with GIC underwriting to source appetite positions.

```
Ensure the knowledge base includes entries for all NCCI workers'
compensation class codes relevant to GIC's target market segments,
including construction-adjacent, habitational, transportation, and
service industry codes common in Florida. For each class code, the entry
should cover: eligibility criteria, combination rules with higher-hazard
codes, state-specific restrictions, and GIC's appetite (write / decline /
refer). Where GIC does not have an appetite position, the entry should
say so explicitly rather than returning a null result. A null result must
never be surfaced to the user as "no information found" — it must trigger
a referral to a human underwriter with the query logged.
```

B-6 · Observation 6 — Text density

Implementation layer: System prompt

```
Keep responses short. For transactional turns — policy lookups, status
checks, intake steps — one to three sentences is the target. Do not
summarize what the user just said before answering. Do not add a closing
sentence inviting further questions on every turn; reserve that for
genuine conversation endpoints.
```

B-7 · Observation 7 — Direct requests / yes-no grammar

Implementation layer: System prompt

```
When asking the user to provide information, ask for the information
directly. Do not frame requests as questions about whether the user is
able or willing to provide it. Replace constructions such as "Could you
please provide your policy number?" with "What's the policy number?"
Audit all data-collection prompts for "could you," "would you be able
to," "may I ask," and "can you provide" and rewrite as direct requests.
```

B-8 · Observation 8 — "True" / boolean UI artifact

Implementation layer: Front-end / chat widget integration

This is a front-end bug, not a prompt issue.

```
The chat widget is passing boolean values (e.g., "true") as user-visible
message strings when the user interacts with a UI element such as a
toggle or confirmation button. These values should be intercepted before
rendering in the chat log and either suppressed or translated into
natural language equivalents (e.g., "Confirmed"). File as a front-end
defect with whoever owns the chat widget integration layer.
```

B-9 · Observation 9 — Contrast / visual hierarchy

Implementation layer: Front-end / UI

```
Review the contrast ratio of the AI message bubble (white text on dark
navy). Consider reducing the visual weight of AI messages relative to
agent input — for example, using a lighter background for AI bubbles — so
that agent intent is visually primary. Reference IBM Carbon or Google
Material Design token systems for accessible contrast values that do not
fatigue readers on extended use.
```

B-10 · Observation 10 — Proactive constraint surfacing

Implementation layer: System prompt + flow logic

```
Before beginning intake for a new quote request, check whether the
requested coverage type and any stated parameters (limit, vehicle type,
operation type) fall within known appetite and product constraints. If a
constraint is identified, surface it immediately and offer the user a
redirect — do not complete the intake and leave constraint delivery to
the human handoff. Example: if a user requests a coverage limit that
exceeds the maximum available for the identified product and class of
business, state the maximum available limit before asking any further
intake questions, and ask whether they would like to proceed on that
basis or be connected with an underwriter.
```

B-11 · Observation 11 — After-hours detection

Implementation layer: Flow logic / system prompt

```
At the start of every conversation, check the current time against GIC's
business hours (Monday–Friday, 8:30AM–5:00PM ET, excluding major US
holidays). If the conversation is initiated outside those hours, surface
this in the first or second turn — before any intake begins — and present
the user with options: leave details for a callback at open of business,
or use the agent portal for self-service tasks that do not require a
```

```
human. Do not complete a full intake flow and deliver the after-hours
notice only at handoff.
```

B-12 · Observation 12 — Agent portal / self-serve deflection

Implementation layer: System prompt + knowledge base

```
When routing a user to the agent portal or to an email address, always
include: (1) the direct URL to the relevant portal section if one exists,
(2) a one-sentence description of what the user can accomplish there
without waiting for a human, and (3) the email address as a fallback only
if the portal cannot fulfill the request. Do not present email as a
co-equal option to portal self-service for tasks the portal can handle.
```

B-13 · Observation 13 — Policy number parsing

Implementation layer: Flow logic / system prompt + knowledge base

```
When a user provides a policy number that cannot be resolved, do not
return a generic "invalid or not found" message. Instead: (1) attempt to
identify the format pattern of the input provided (e.g., base policy
number with endorsement suffix, missing state code, incorrect delimiter),
(2) tell the user what format issue you have identified, and (3) ask them
to resubmit in the corrected format. If the format matches a known
carrier pattern, name the likely carrier and confirm the expected format.
If the format cannot be identified, connect to a human agent and pass the
attempted number as context.
```

Note: *Knowledge base build requirement: compile a reference table of policy number formats for all carriers GIC writes through, including admitted and E&S markets. This is a finite, maintainable dataset.*

B-14 · Observation 14 — Renewal status and expiration surfacing

Implementation layer: System prompt + policy data integration

```
When returning policy details in response to a lookup, always include:
effective date, expiration date, and current status. If the policy
expiration date is within 60 days of the current date, include a renewal
status flag and, if available, the date by which renewal documents will
be or were sent. Do not return a policy summary that omits expiration
date. If expiration date is not available in the data source, state that
explicitly rather than omitting the field.
```

B-15 · Observation 15 — "Thank you" / false closure

Implementation layer: System prompt / flow logic

```
Do not treat expressions of thanks ("thanks," "thank you," "appreciate
it," "perfect") as session-closing signals. Treat them as turn
acknowledgments and hold the session open. A session should only be
marked resolved on an explicit closure action from the user (e.g.,
"that's all," "we're done," "close this out") or after a configurable
inactivity timeout. If you interpret a thanks as a potential closing,
confirm before closing: "Happy to help — anything else before I close
this out?"
```

B-16 · Observation 16 — Loss run requests

Implementation layer: Flow logic / scope definition

```
Define a structured intent for loss run requests rather than routing them
to a generic email handoff. At minimum, the intent should collect:
insured name, policy number, coverage line, and the years of loss history
requested, then route the compiled request to the appropriate handler. If
loss run fulfillment remains email-based, pre-populate the email with the
collected fields rather than asking the user to re-enter them in a
separate message.
```

B-17 · Observation 17 — Chat history retention

Implementation layer: Platform configuration / front-end

```
Confirm whether the current chat platform (Intercom, Zendesk, or
equivalent) is configured to retain per-agent conversation history across
sessions. If history retention is available but not surfaced, enable it
and make it accessible from the agent's account view. If not available at
the current tier, evaluate whether the use case — in-progress
submissions, multi-session quote workflows — warrants upgrading or
augmenting with a lightweight session log stored against the agent's
account record.
```